

Privacy Preserving Speech Processing

Manas A. Pathak, Bhiksha Raj, Shantanu Rane, Paris Smaragdis

I. INTRODUCTION

Speech is one of the most private forms of communication. People do not like to be eavesdropped on. They will frequently even object to being recorded; in fact in many places it is illegal to record people speaking in public, even when it is acceptable to capture their images on video [1]. Yet, when a person uses a speech-based service such as a voice authentication system or a speech recognition service, they must grant the service complete access to their voice recordings. This exposes the user to abuse, with security, privacy and economic implications. For instance, the service could extract information such as gender, ethnicity, and even the emotional state of the user from the recording – factors not intended to be exposed by the user – and use them for undesired purposes. The recordings may be edited to create *fake* recordings that the user never spoke, or to impersonate them for other services. Even derivatives from the voice are risky to expose. E.g. a voice-authentication service could make unauthorized use of the models or voice prints it has for users to try to identify their presence in other media such as YouTube.

Privacy concerns also arise in other situations. For instance, a doctor cannot just transmit a dictated medical record to a generic voice-recognition service for fear of violating HIPAA requirements; the service provider requires various clearances first. Surveillance agencies must have access to all recordings by all callers on a telephone line, just to determine if a specific person of interest has spoken over that line. Thus, in searching for Jack Terrorist, they also end up being able to listen to and thereby violate the privacy of John and Jane Doe.

The need to protect the privacy of users and their data is well recognized in other domains [2]. Any private information that can be gleaned by inspecting a user's interaction with a system must be protected from prying eyes. To this end, techniques have been proposed in the literature for protecting user privacy in a variety of applications including e-voting, information retrieval and biometrics. Yet, the privacy of voice has not been addressed until recently, and the issue of privacy of speech has been dealt with primarily as a policy problem [3], [4], and not as a technology challenge. In this article, we describe recent developments in *privacy-preserving*

frameworks for voice processing. Here, we refer chiefly to secure pattern-matching applications of speech such as speech biometrics and speech recognition, rather than secure speech *communication* techniques, which have already been studied [5].

The goal of the described frameworks is to enable voice-processing tasks in a manner which ensures that no party, including the user, the system, or a snooper, can derive undesired or unintended information from the transaction. This would imply, for instance, that a user may enroll at a voice-authentication system without fear that an intruder *or even the system itself* could capture and abuse his voice or statistical models derived from it. A surveillance agency could now determine if a crime suspect is on a telephone line, but would learn nothing if the speaker is an innocent person. Private speech data may be mined by third parties without exposing the recordings.

These frameworks follow two distinct paradigms. In the first paradigm [6], [7], [8], [9], [10], which we will refer to as the “cryptographic” paradigm, conventional voice-processing algorithms are rendered secure by computing them through secure operations. By recasting the computations as a pipeline of “primitives”, each of which can be computed securely through a combination of *homomorphic encryption* [11], [12], [13], [14], *secure multiparty computation* (SMC) [15] and *oblivious transfer* [16], [17], we ensure that no undesired information is leaked by any party. In this paradigm, the accuracy of the basic voice-processing algorithm can remain essentially unchanged with respect to the original non-private version. The privacy requirements, however, introduce a computational and communication overhead as the computation requires user participation.

The second paradigm modifies voice-pattern classification tasks into a string-comparison operation [18], [10]. Using a combination of appropriate data representation followed by *locality sensitive hashing* schemes both the data to be matched and the patterns they must match are converted to collections of bit strings, and pattern classification is performed by counting exact matches. The computational overhead of this string-comparison framework is minimal compared to the cryptographic framework. Moreover, the entire setup is non-interactive and secure, in the sense that no party learns anything

undesired regardless of how they manipulate the data. This comes at the price of a slight loss of performance, since the modified classification mechanism is not as effective as conventional classification schemes.

In the rest of this paper, we first present a brief primer on speech in Section II and discuss its privacy issues in Section III. We describe the cryptographic approach in Section IV and the string-comparison approach in Section V. Finally in Section VI we present our conclusions. In our discussion, we will assume a *user* who possesses voice data and a *system* who must process it. This nomenclature may not always be appropriate; nevertheless we will retain it for consistency.

II. A BRIEF PRIMER ON SPEECH PROCESSING

The speech processing applications we consider all deal with making inferences about the information in the recorded signal. *Biometric* applications attempt to determine or confirm the identity of the speaker of a recording. *Recognition* applications attempt to infer *what* was spoken in the recording. We present below a very brief description of the salient aspects of these applications as they pertain to this paper. The description is not intended to be exhaustive; for more detailed information we refer readers to the various books and papers on the topics, *e.g.* [19].

In all cases, the problem of inference is treated as one of statistical pattern classification. Classification is usually performed through a Bayes classifier. Let \mathcal{C} be a set of candidate classes that a recording \mathbf{X} might belong to. Let $P(\mathbf{X}|\mathcal{C})$ be the probability distribution of speech recordings \mathbf{X} from class \mathcal{C} . $P(\mathbf{X}|\mathcal{C})$ is usually represented through a parametric model, *i.e.* $P(\mathbf{X}|\mathcal{C}) \approx P(\mathbf{X}; \lambda_{\mathcal{C}})$ where $\lambda_{\mathcal{C}}$ are the parameters of the class \mathcal{C} and are learned from data. Classification is performed as

$$\hat{\mathcal{C}} = \arg \max_{\mathcal{C} \in \mathcal{C}} \log P(\mathbf{X}; \lambda_{\mathcal{C}}) + \log P(\mathcal{C}) \quad (1)$$

where $P(\mathcal{C})$ represents the *a priori* bias for class \mathcal{C} . Stated thus, the primary difference among the applications considered lies in the definition of the candidate classes in \mathcal{C} and the nature of the model $P(\mathbf{X}; \lambda_{\mathcal{C}})$.

Before proceeding, we note that the classifiers do not work directly from the speech signal. Instead, the signal is converted to a sequence of *feature vectors*, typically “Mel-frequency cepstral coefficients” (MFCC) [20]. To generate these, the signal is segmented into overlapping “frames”, each typically 25ms wide, with an overlap of 15ms between adjacent frames. From each frame, a vector of MFCC (or similar) features is derived, which may be further augmented with their temporal

differences and double-differences. For our purposes, it suffices to know that when we refer to a speech recording we actually refer to the sequence of feature vectors $\mathbf{X} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T]$. For the privacy-preserving frameworks described later, we will assume that the user’s (client) device can compute these features. The conversion of signals to features by itself provides no privacy, since they can be used for pattern classification, and can even be inverted to generate intelligible speech signals [21].

A. Biometric Applications: Identification and Authentication

Biometric applications deal with determining the identity of the speaker. Here, the set \mathcal{C} in (1) is the set of candidate speakers who may have spoken in a recording. In addition, a “universal” speaker U , representing the aggregate of all speakers not otherwise in \mathcal{C} is included in the set. U may be viewed as the “none-of-the-above” option – classifying a recording as U is equivalent to stating that the speaker is unknown. In speaker *identification* systems \mathcal{C} comprises a collection of speakers S (possibly including U) for whom models λ_S are available. In speaker *authentication*, \mathcal{C} comprises only the speaker who has claimed to have spoken in the recording, and the universal speaker U . $P(\mathcal{C})$ in (1) now becomes a tuning parameter to bias the classification towards specific speakers or U .

Typically, the individual feature vectors in any recording \mathbf{X} are assumed to be IID and distributed according to a Gaussian mixture model (GMM) [22]. Thus, for any speaker S in \mathcal{C} , $P(\mathbf{X}; \lambda_S)$ is assumed to have the form $P(\mathbf{X}; \lambda_S) = \prod_t P(\mathbf{x}_t; \lambda_S)$, where

$$P(\mathbf{x}_t; \lambda_S) = \sum_{k=1}^K w_k^S \mathcal{N}(\mathbf{x}_t; \mu_k^S, \Sigma_k^S). \quad (2)$$

Here K is total number of Gaussians in $P(\mathbf{x}_t; \lambda_S)$, $\mathcal{N}()$ represents a Gaussian density, and w_k^S , μ_k^S and Σ_k^S are the mixture weight, mean vector and covariance matrix of the k^{th} Gaussian in the mixture. The parameters of the GMM are expressed as $\lambda_S = \{w_k^S, \mu_k^S, \Sigma_k^S \forall k = 1, \dots, K\}$.

λ_U , the parameters of the GMM for “universal speaker” U are learned from a large collection of speech recordings from many speakers. λ_U is often called a “Universal Background Model”, or UBM. The GMM parameters λ_S for any speaker S are learned from a collection of recordings for the speaker, using the EM algorithm. When the amount of training data for the speaker are limited, *e.g.* from enrollment recordings in an authentication system, and insufficient to learn GMM

parameters robustly, they are learned by *adapting* the parameters of the UBM to the data from the speaker using an MAP estimation algorithm [23], [24]. Classification of speakers with the GMMs now proceeds as

$$\hat{S} = \operatorname{argmax}_{S \in \mathcal{C}} \sum_{t=1}^T \log P(\mathbf{x}_t; \lambda_S) + \theta_S. \quad (3)$$

where θ_S encodes our prior biases as mentioned earlier.

B. Recognition Applications

In speech *recognition* systems, \mathcal{C} in (1) represents the collection of all possible word sequences that a person may say in the recording. *Phrase spotting* systems only include the specific phrases of interest, along with a background model, usually called a “garbage model” in this context, which represents the “none of the above option”. *Isolated word recognition* systems assume that only a single word was spoken in the recorded segment being analyzed; \mathcal{C} here comprises the vocabulary of words to recognize. In *continuous speech recognition* (CSR) systems, \mathcal{C} represents the set of all possible sentences a person may speak. This set can be very large, and even infinite in size. To make classification manageable, the set of all possible sentences is represented as a compact, loopy word graph, which conforms to a *grammar* or *n*-gram language model that embodies the *a priori* probabilities $P(C)$ [19].

The probability distribution $P(\mathbf{X}; \lambda_C)$ for each class C is usually modelled by a hidden Markov model (HMM). The theory of HMMs is well known; we only reiterate the salient aspects of it here [25], [19]. An HMM is a model for time-varying processes and is characterized by a set of states $[s_1, \dots, s_M]$ and an associated set of probability distributions. According to the model, the process transitions through the states according to a Markov chain. After each transition it draws an observation vector from a probability distribution associated with its current state. The parameters characterizing the HMM for any class C are i) the initial state probabilities $\Pi^C = \{\pi_i^C, i = 1, \dots, M\}$, where π_i^C represents the probability that at the first instant the process will be in state s_i ; ii) transition probabilities $\mathbf{A}^C = \{a_{i,j}^C, i = 1, \dots, M, j = 1, \dots, M\}$ which represent the probability that, given the process is in state s_i at any time, it will jump to s_j at the next transition, and iii) the set of *state output probability distributions* $\{P^C(\mathbf{x}_t; \Lambda_i^C)\}$ associated with each state. In speech recognition systems the $P^C(\mathbf{x}_t; \Lambda_i^C)$ are generally modeled as Gaussian mixture densities: $P^C(\mathbf{x}_t; \Lambda_i^C) = \sum_k w_{i,k}^C \mathcal{N}(\mathbf{x}_t; \mu_{i,k}^C, \Sigma_{i,k}^C)$. Thus the parameters for class

C are $\lambda_C = \{\Pi^C, \mathbf{A}^C, \mathbf{\Lambda}^C\}$, where $\mathbf{\Lambda}^C = \{\Lambda_i^C \forall i = 1, \dots, M\}$ and $\Lambda_i^C = \{w_{i,k}^C, \mu_{i,k}^C, \Sigma_{i,k}^C \forall k\}$.

To compute $P(\mathbf{X}; \lambda_C)$ for any class C , we must employ the following recursion, commonly known as the *forward* recursion. Here, the term $\alpha^C(t, i)$ represents the total probability that the process arrives at state i after t transitions and generates the partial sequence $\mathbf{x}_1, \dots, \mathbf{x}_t$, i.e., $\alpha^C(t, i) = P(\mathbf{x}_1, \dots, \mathbf{x}_t, \text{state}(t) = i; \lambda_C)$.

$$\begin{aligned} \alpha^C(1, i) &= \pi_i^C P^C(\mathbf{x}_1 | i) \\ \alpha^C(t, i) &= P(\mathbf{x}_t | i, C) \sum_j \alpha^C(t-1, j) a_{j,i}^C \quad \forall t > 1 \\ P(\mathbf{X}; \lambda_C) &= \sum_i \alpha^C(T, i) \end{aligned} \quad (4)$$

where T is the total number of feature vectors in \mathbf{X} .

$P(\mathbf{X}; \lambda_C)$ computed in the above manner considers *all possible* state sequences that the process may have followed to generate \mathbf{X} . It can be used in (1) for phrase spotting and isolated-word recognizers. In continuous speech recognition, however, the classes are word sequences, which are collapsed into a compact word graph [26] and computing $P(\mathbf{X} | C)$ through (4) is not feasible for individual word sequences. Here, $P(\mathbf{X}; \lambda_C)$ is replaced by the probability of the *most likely* state sequence though the HMM for C , and the classification is performed as

$$\hat{C} = \operatorname{argmax}_C \log P(C) + \max_{\mathbf{s}} \log P(\mathbf{X}, \mathbf{s}; \lambda_C) \quad (5)$$

where \mathbf{s} is the state sequence followed by the process. Unlike classification based on forward probabilities, this estimation can be performed over the set of word sequences represented by a word graph: the word graph is *composed* into a large HMM by replacing each edge in the graph by the HMM for the word it represents [26]. The word sequence corresponding to the most probable state sequence through the resulting HMM is guaranteed to be identical to the one obtained by (5). Even in contexts outside continuous speech recognition, (5) is frequently used as a generic substitute for (1), since it is more efficient to compute.

The term $\max_{\mathbf{s}} \log P(\mathbf{X}, \mathbf{s} | C)$ in turn can be computed very efficiently using a dynamic programming algorithm known as the *Viterbi* algorithm, which implements the following recursion. Here $\Gamma^C(t, i)$ represents the log of the joint probability of $\mathbf{x}_1 \dots \mathbf{x}_t$ and the most

probable state sequence that arrives at state s_i at time t .

$$\begin{aligned}\Gamma^C(1, i) &= \log \pi_i^C + \log P(\mathbf{x}_1; \Lambda_i^C) \\ \delta^C(t, i) &= \underset{j}{\operatorname{argmax}} \Gamma^C(t-1, j) + \log a_{j,i}^C \quad \forall t > 1 \\ \Gamma^C(t, i) &= \log P(\mathbf{x}_t; \Lambda_i^C) + \Gamma^C(t-1, \delta^C(t, i)) \\ &\quad + \log a_{\delta^C(t, i), i}^C \\ \max_{\mathbf{s}} \log P(\mathbf{X}, \mathbf{s} | C) &= \max_i \Gamma^C(T, i)\end{aligned}\quad (6)$$

$\delta^C(t, i)$ is a table of “backpointers” from which the most likely state sequence $\operatorname{argmax}_{\mathbf{s}} \log P(\mathbf{X}, \mathbf{s}; \lambda_C)$ can be determined by tracing backwards recursively as $s_T = \operatorname{argmax}_i \Gamma^C(T, i)$, $s_t = \delta^C(t+1, s_{t+1})$, where s_t is the t^{th} state in optimal state sequence. In continuous speech recognition systems, in particular, the backtraced optimal state sequence traces a path through the word graph and is used to determine the spoken word sequence.

III. PRIVACY ISSUES IN SPEECH PROCESSING

From the discussion of the previous section, the key component of all the above applications is the computation of the *class score* $\log P(\mathbf{X}; \lambda_C)$. In conventional — non-private — implementations of these applications, the system providing the application has access to data \mathbf{X} , which represents the user’s voice. This enables it to manipulate or misuse the data, raising privacy concerns. The alternative, *i.e.*, permitting the user to access the models λ_C , is equally unacceptable. In many situations, the models are the system’s intellectual property. Furthermore, in voice-data mining situations, the system’s models also represent the patterns it is searching for; revealing these to external parties may not be acceptable. A voice-based authentication system where the user himself has access to the models cannot be considered an effective authenticator.

Thus, for maximally protecting the user and the system from one another, the system must not learn the user’s data \mathbf{X} , while the user must not be able to infer the system’s models λ_C . For the speech processing applications discussed above, this means that log probability terms computed from GMMs, forward probabilities, and best-state-sequence probabilities (also called Viterbi scores) must all be computed without revealing \mathbf{X} to the system or λ_C to the user. An additional twist arises in the case of *authentication* systems, where the system must be prevented from making unauthorized use of the models it has for the user. In this case, the model λ_S must itself be in a form that can only be useful when the system engages with the appropriate user.

The above deals with the computation of scores for any class. But what about the *final* outcome of the

classification? This, too, has an intended recipient. For instance, in most biometric applications it is the *system* that must receive the outcome of the classification; however for recognition systems the user obtains the result. Thus we also stipulate that the outcome is only revealed to the intended recipient. We refer to any computational mechanisms that enable the above requirements as *private* computation, as opposed to conventional non-private methods that make no guarantees to privacy. The next two sections describe two frameworks that enable such private processing.

IV. THE CRYPTOGRAPHIC APPROACH

The cryptographic approach treats private speech processing as an instance of *secure two-party computation* [15]. Consider the case where two parties, Alice and Bob have private data a and b respectively and they want to compute the result of a function $f(a, b)$. Any computational protocol to calculate $f(a, b)$ is said to be *secure* only if it leaks no more information about a and b than what either party can gain from learning the result c . We assume a *semi-honest model* for the parties where each party follows the protocol but could save messages and intermediate results to learn more about other’s private data. In other words, the parties are honest but curious and will follow the agreed-upon protocol but will try to learn as much as possible from the data flow between the parties. We return to the issue of honesty later in this section.

We recast the conventional speech processing algorithms described in Section II as a pipeline of primitive computations, and show how to execute each primitive in a *computationally secure* manner, *i.e.*, a computationally bounded participant should not be able to derive private information possessed by the other participant from the computation. The output of each stage of the pipeline is either distributed across both participants in the form of random additive shares, or arrives at one of the participants in a *locked* form, where the other participant holds the key. Thus both participants must interact to perform the computations until the final outcome of the overall computation arrives at the correct participant.

A. Secure Primitives for Speech Processing

To keep the discussion simple, we will consider that there are only two parties, Alice and Bob, engaged in the computation. The reader may consider Alice as the client or end-user who possesses a speech signal to be analyzed while keeping it private from Bob, the remote system which has the model parameters, some or all of which must be kept private from Alice.

To enable private computation, we will utilize public-key *homomorphic* encryption schemes, which enable operations to be performed on encrypted data [11]. We keep the development simple by considering an additively homomorphic cryptosystem [12], [13], [14] with encryption function $E[\cdot]$. Such a cryptosystem satisfies $E[x] \cdot E[y] = E[x + y]$ and $(E[x])^y = E[xy]$ for integer messages x, y . For a discussion of other flavors of homomorphic cryptosystems, namely multiplicatively homomorphic, 2-DNF homomorphic and fully homomorphic cryptosystems, the reader is referred to a companion article in this issue [27].

We assume that the cryptosystem is semantically secure [28], i.e., by using fresh random parameters during encryption but not during decryption, a given plaintext may be mapped to different ciphertexts every time encryption is performed. This makes the cryptosystem, and hence the protocols discussed, resilient to a chosen plaintext attack (CPA). We further assume that while Alice and Bob share a public encryption key, Alice is the only person with a decryption key that reverses $E[\cdot]$. This means that the system may encrypt data if needed, but only the end-user may decrypt it. However, if required the situation can be reversed by mirroring the protocols, with relatively minor additional changes. Assume that Alice and Bob own n -length integer vectors \mathbf{x} and \mathbf{y} respectively. In what follows, $E[\mathbf{x}]$ denotes a vector containing encryptions of the individual elements of \mathbf{x} , i.e., $(E[x_1], E[x_2], \dots, E[x_n])$.

Additive Secret Sharing

Bob has an encrypted value $E[x]$. He wishes to share it as random additive clear-text shares with Alice. He chooses an integer b at random, and sends $E[x] \cdot E[-b] = E[x - b]$ to Alice. Bob retains b as his additive share. Alice decrypts $x - b = a$ which is her additive share. We will represent this operation in short-hand by saying Alice and Bob receive a and b such that $a + b = \text{SHARE}(x)$.

Secure Inner Product (SIP)

Alice and Bob want to compute uninformative additive shares a, b respectively of the inner product $\mathbf{x}^\top \mathbf{y}$. There are several ways to achieve this [29], [30], [31], but for expository purposes, we focus on a simple approach using additively homomorphic functions [32]. Alice sends element-wise encryptions $E[x_i]$ to Bob. Bob computes $\prod_{i=1}^n E[x_i]^{y_i} = E[\sum_{i=1}^n x_i y_i] = E[\mathbf{x}^\top \mathbf{y}]$. Alice and Bob can then receive additive shares a and b of $E[\mathbf{x}^\top \mathbf{y}]$ as $a + b = \text{SHARE}(\mathbf{x}^\top \mathbf{y})$. Observe that Bob operates in the encrypted domain and cannot discover \mathbf{x} . Similarly,

Alice does not know b , and hence cannot discover \mathbf{y} .

When Bob possesses only $E[\mathbf{y}]$ rather than \mathbf{y} as assumed above, it is still possible — using a trick involving additive secret sharing — for Alice and Bob to obtain additive shares of $\mathbf{x}^\top \mathbf{y}$. As a notational shorthand, when we invoke any variant of this protocol, we will just state that Alice and Bob obtain additive shares a, b , such that $a + b = \text{SIP}(\mathbf{x}, \mathbf{y})$.

Secure Logsum (SLOG)

Suppose that $\mathbf{x} + \mathbf{y} = (\ln z_1, \ln z_2, \dots, \ln z_n)$. By a slight abuse of notation, denote the vector of the elementwise logarithms and elementwise exponents of the elements of \mathbf{x} as $\ln \mathbf{x}$ and $e^{\mathbf{x}}$ respectively. Alice and Bob wish to obtain uninformative additive shares, a and b such that $a + b = \ln(\sum_{i=1}^n z_i)$, which is the “logsum” operation that gives the protocol its name. We note that $\sum_{i=1}^n z_i = \sum_{i=1}^n e^{x_i + y_i}$ and achieve the desired secret sharing using the following protocol [7]:

- 1) Alice chooses a at random. Then Alice and Bob compute additive shares q, s such that $q + s = \text{SIP}(e^{x-a}, e^y)$ using the SIP protocol above. Bob combines these shares to obtain the inner product ϕ .
- 2) Bob computes $b = \ln \phi = -a + \ln(\sum_{i=1}^n e^{x_i + y_i}) = -a + \ln(\sum_{i=1}^n z_i)$, which gives the desired result.

In the first step above, Alice and Bob employ additive secret sharing in the exponent, which is equivalent to multiplicative secret sharing. The parameter a should be chosen large enough because multiplicative secret sharing is not as secure as standard additive secret sharing. We present this protocol to illuminate the fact that homomorphic functions can be manipulated to compute useful, non-obvious functions. The same functionality can be obtained in a secure manner using other cryptographic primitives, e.g., garbled circuits. To refer to this protocol henceforth, we will state that Alice and Bob obtain additive shares a, b , such that $a + b = \text{SLOG}(\ln \mathbf{z})$.

In an alternate scenario, Bob possesses $E[\ln z_1], E[\ln z_2], \dots$. He wishes to obtain $E[\log \sum_i z_i]$. He uses the *SHARE* protocol to share $E[\ln z_i] \forall i$ with Alice and proceeds as earlier. Finally, Alice sends $E[a]$ to Bob who computes $E[\log \sum_i z_i] = E[a]E[b]$. Note that in the process Alice and Bob also obtain additive shares of the outcome. We will denote this operation by stating that Bob receives c such that $c = \text{SLOG}(\ln \mathbf{z})$.

Secure Maximum Index (SMI)

Alice and Bob wish to compute additive shares, such that $a + b = \arg\max_i x_i + y_i$ without revealing their data to each other. This is achieved by exploiting homomorphic encryption to perform blind-and-permute

operations [33]. Let $\mathbf{z} = \mathbf{x} + \mathbf{y}$. Then, the goal is to privately compute additive shares of the index of the largest element of \mathbf{z} . To accomplish this, Alice chooses a secret permutation π_A on the index set $\{1, 2, \dots, n\}$. Similarly Bob chooses a secret permutation π_B on the same set. The outcome of the blind-and-permute protocol is that Alice and Bob each obtain additive shares of $\pi_A \pi_B \mathbf{z}$. Neither can reverse the other person's permutations. However, given the permuted shares, Alice and Bob can use repeated instantiations of the millionaire protocol [17], obtain the index of the maximum element in the permutation of \mathbf{z} , from which, using their respective permutations, they can obtain shares of the index of the maximum element in \mathbf{z} . To see the steps involved in minimum finding based on homomorphic encryption, refer to a companion article submitted to this special issue [34]. Alternatively the entire minimum finding protocol may be executed using garbled circuits [35]. As a shorthand notation, we say that Alice and Bob obtain additive shares a, b , such that $a + b = SMI(\mathbf{x}, \mathbf{y})$.

In an alternate scenario, Bob has two encrypted numbers $E[x]$ and $E[y]$ and desires to find which of the two is larger, he can engage in protocols such as those described in [36], which employ *threshold encryption* or *secret sharing* schemes in which both parties must cooperate to perform decryption such that one of them can obtain the answer without exposing x or y to either. We denote this also as $\max(x, y) = SMI(x, y)$; this should not result in confusion since the actual operation used will be clear from the context.

Secure Maximum Value (SMV)

Alice and Bob wish to compute uninformative additive shares a, b such that $a + b = \max_i x_i + y_i$. This is achievable using a protocol similar to the one described above, with the only difference being in the last step, which – instead of revealing the index of the maximum – reveals additive shares of the maximum value. To invoke this protocol, we will state that Alice and Bob obtain additive shares a, b , such that $a + b = SMV(\mathbf{x}, \mathbf{y})$ [7].

Other similar protocols may be defined. In particular, it is often essential to compute distance measures, such as the Hamming or Euclidean or Absolute distance between vectors \mathbf{x} and \mathbf{y} held privately by Alice and Bob respectively. Protocols for these computations can be efficiently designed using homomorphic functions and are an integral part of privacy-preserving nearest neighbor methods which are covered in a separate article in this issue [34].

Practical Considerations

An important practical consideration is that we cannot

always assume that \mathbf{x} and \mathbf{y} are integer vectors. Indeed, speech processing routinely uses probabilistic models, such as Gaussian Mixture Models (GMMs) and Hidden Markov Models (HMMs) in which the model parameters are floating point values between 0 and 1. A particular problem that is nearly ubiquitous in the iterative algorithms used to perform modern speech processing is that multiplication of probability values result in extremely small numbers. One way to mitigate this issue of floating point precision is to take the logarithm of the probability values and operate exclusively in the logarithmic domain. In fact, the need to perform operations on the logarithms of the probability values is what makes the SLOG protocol so useful. A second way to mitigate the issue of extremely small probability values is to appropriately scale the probability values by a large constant, e.g., 10^6 prior to all encryptions and compensate for the scaling after decryption.

B. Computing Scores Privately for Speech Processing

The computations involved in speech processing tasks can now be cast in terms of the above primitives. We first consider how the various required scores can be computed privately.

The Gaussian as a Dot Product

The fundamental component of speech processing models is the Gaussian, since observation distributions are generally modeled as Gaussian mixtures. The form of the log of a multi-variate Gaussian is well known:

$$\log P(\mathbf{x}) = -0.5(\mathbf{x} - \mu)^T \Sigma^{-1}(\mathbf{x} - \mu) - 0.5 \log(2\pi)^D |\Sigma|$$

where D is the dimensionality of the data, and μ and Σ are the mean vector and covariance matrix of the Gaussian. It is fairly simple to show that this can be manipulated into the form $\tilde{x}^T \tilde{W} \tilde{\mathbf{x}}$ [7] where $\tilde{\mathbf{x}}$ is obtained by extending \mathbf{x} as $\tilde{\mathbf{x}} = [\mathbf{x}^T 1]^T$, and

$$\tilde{W} = \begin{bmatrix} -0.5\Sigma^{-1} & \Sigma^{-1}\mu \\ 0 & w^* \end{bmatrix} \quad (7)$$

where $w^* = -0.5\mu^T \Sigma^{-1} \mu - 0.5(2\pi)^D \log |\Sigma| + \log \text{prior}$, and *prior* captures an *a priori* probability associated with the Gaussian. In the case of a solitary Gaussian, *prior* = 1; however if the Gaussian is one from a mixture, *prior* represents the mixture weight for the Gaussian. We can reduce the above computation further to a single inner product $\tilde{\mathbf{x}}^T \tilde{W}$, where $\tilde{\mathbf{x}}$ is a *quadratically extended* feature vector derived from $\tilde{\mathbf{x}}$ which consists of all pairwise product terms $\tilde{x}_i \tilde{x}_j$ of all components $\tilde{x}_i, \tilde{x}_j \in \tilde{\mathbf{x}}$, and \tilde{W} is obtained

by unrolling \tilde{W} into a vector. In this representation $\log \mathcal{N}(x; \mu, \Sigma) = \bar{\mathbf{x}}^\top W$.

Privacy preserving computation of a Log Gaussian

Input: Alice possesses a vector \mathbf{x} . Bob possesses a Gaussian parameterized by $\lambda = \{\mu, \Sigma\}$.

Output: Alice and Bob obtain random additive shares a and b such that $a + b = \log P(\mathbf{x}; \lambda)$.

- 1) Alice computes the extended vector $\bar{\mathbf{x}}$ from \mathbf{x} . Bob arranges his model into a vector W as explained above.
- 2) Alice and Bob engage in the *SIP* protocol to obtain additive shares a and b : $a + b = \text{SIP}(\bar{\mathbf{x}}, W)$.

Note that this is possible even if Bob only possesses an *encrypted* version of W . Note also that given only $E[\mu]$, $E[\Sigma]$ and $E[\text{prior}]$ Bob can obtain $E[W]$ using the protocol given in [9]. In either case, Alice and Bob obtain no information about each other's data.

Computing the Logarithm of a Gaussian Mixture Privately

Input: Alice possesses \mathbf{x} . Bob possesses $\Lambda = \{\lambda_1, \lambda_2, \dots, \lambda_K\}$, the parameters of a Gaussian mixture with K Gaussians, each with parameter λ_i .

Output: Alice and Bob receive a and b such that $a + b = \log P(\mathbf{x}; \Lambda)$.

- 1) Bob arranges the parameters λ_i of each Gaussian into a vector W_i , $i = 1, \dots, K$.
- 2) For each $i = 1, \dots, K$, Alice and Bob engage in the *SIP* protocol to obtain additive shares c_i and d_i such that $c_i + d_i = \text{SIP}(\bar{\mathbf{x}}, W_i)$. Note that if W_i is available in unencrypted form, Alice only needs to transmit $E[\bar{\mathbf{x}}]$ once for the entire Gaussian mixture.
- 3) Alice and Bob apply the *SLOG* protocol to obtain $a + b = \text{SLOG}(\ln \mathbf{g})$, where $\mathbf{g} = [c_1 + d_1, \dots, c_K + d_K]$, where $c_i + d_i$ is the log of the i^{th} Gaussian.

As before, Alice and Bob do not learn about each other's data. We will refer to this operation as $a + b = \text{SMOG}(\mathbf{x}, \Lambda)$.

Computing the logarithm of an HMM forward score privately

Input: Alice possesses $\mathbf{X} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T]$. Bob possesses an M -state HMM $\Gamma = \{\Pi, \mathbf{A}, \Lambda\}$, with initial state probabilities $\Pi = [\pi_1, \dots, \pi_M]$, a transition matrix \mathbf{A} comprising vectors $\mathbf{a}_i = [a_{i,1}, \dots, a_{i,M}]$ and a set of state output Gaussian mixture densities $\Lambda = \{\Lambda_1, \Lambda_2, \dots, \Lambda_M\}$.

Output: Alice and Bob obtain a and b such that $a + b = \log P(\mathbf{X}; \Gamma)$.

Step I State output density computation

- 1) For all $t = 1 \dots T$, $i = 1 \dots M$
 - a) Alice and Bob engage in the *SMOG* protocol to obtain additive shares $g_{t,i} + h_{t,i} = \text{SMOG}(\mathbf{x}_t, \Lambda_i)$. Alice sends $E[g_{t,i}]$ to Bob.
 - b) Bob computes $q_{t,i} = E[g_{t,i}]E[h_{t,i}] = E[g_{t,i} + h_{t,i}] = E[\log P(\mathbf{x}_t; \Lambda_i)]$ to obtain the encrypted value of the logarithm of the state output distribution for state i on \mathbf{x}_t .

In the protocol below we use the notation $\alpha_l(t, i) = \log \alpha(t, i)$ and $\alpha_l(t) = [\log \alpha(t, 1), \dots, \log \alpha(t, M)]$. We represent $\{q_{t,i} \forall i\}$ obtained by Bob in the above computation as \mathbf{q}_t . Now Bob precomputes $E[\log \mathbf{a}_i] \forall i$. The recursions of the forward probability computation in HMMs given by Equation 4 can now be computed in the log domain securely as follows.

Step II Forward Algorithm

- 1) Bob computes $E[\alpha_l(1)] = \mathbf{q}_t \cdot E[\log \pi]$.
- 2) For $t = 2 \dots T$, $i = 1 \dots M$
 - a) Bob engages with Alice in the *SLOG* protocol to obtain $c = \text{SLOG}(\log \mathbf{a}_i + \alpha_l(t-1))$. Note that $c = E[\log \sum_j \alpha(t-1, j) a_{j,i}]$.
 - b) Bob computes $E[\alpha_l(t, i)] = c \cdot q_{t,i}$
- 3) Alice and Bob engage in the *SLOG* protocol to obtain additive shares $a + b = \text{SLOG}(\alpha_l(T))$. Note that a and b are additive shares of $\log P(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T; \Gamma)$.

We will refer to this operation as $a + b = \text{SFWD}(\mathbf{X}, \Gamma)$.

Secure Viterbi Algorithm

The Viterbi algorithm is very similar to the forward algorithm, except for two key differences. First, instead of *adding* all incoming probabilities into any state in (4), we choose the maximum in (6). Additionally, Alice must also receive the optimal state sequence. Ideally, state indices would be permuted in the received sequence; however, for simplicity, we omit the additional complexity involved with permuting the indices. Using the notation of (6): $\delta(t, i)$ refers to the “best” predecessor for state s_i at time t . $\Gamma(t, i)$ is the joint log likelihood of the most probable state sequence arriving at state s_i at time t and the observation sequence $\mathbf{x}_1, \dots, \mathbf{x}_t$. We will use the notation that a bold character represents a vector that aggregates the scalar values represented by the corresponding unbolded symbol for all states. For instance $\Gamma(t) = [\Gamma(t, 1), \Gamma(t, 2), \dots, \Gamma(t, M)]$.

Input: Alice possesses $\mathbf{X} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T]$. Bob possesses an M -state HMM $\Gamma = \{\Pi, \mathbf{A}, \Lambda\}$.

Output: Alice and Bob obtain additive shares a and b such that $a + b = \log \max_{\mathbf{s}} P(\mathbf{X}, \mathbf{s}; \Gamma)$. Alice receives

the optimal state sequence $\bar{s} = s_1, \dots, s_T$.

The protocol uses operations similar to the ones described in the previous section and in the secure forward algorithm described above. For the detailed steps of the protocol, the reader is referred to [7].

C. Implementing Secure Speech Techniques

Having set up the operations as described above, we now consider how they may be applied to our speech problems. We will generally assume that all models have already been learned by the system, and, where required (e.g. speaker authentication), the models are encrypted. To avoid making the development unnecessarily complex, we do not describe how to privately learn GMM/HMM parameters [7], or how to privately adapt an existing background GMM or UBM to a user's enrollment data [9].

Speaker authentication

Input: The system possesses encrypted models Λ_S for the speaker and clear-text models Λ_U as the universal background model. The user possesses vectors $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T$.

Output: The system authenticates the user.

- 1) For each $t = 1 \dots T$
 - a) The user and the system engage in the *SMOG* operation with \mathbf{x}_t and Λ_S to obtain additive shares a_t^S and b_t^S .
 - b) The user and the system engage in the *SMOG* operation with \mathbf{x}_t and Λ_U to obtain additive shares a_t^U and b_t^U .
- 2) The user computes $A^S = \sum_t a_t^S$ and $A^U = \sum_t a_t^U$. The system computes $B^S = \sum_t b_t^S$ and $B^U = \sum_t b_t^U$.
- 3) The user and system engage in an *SMI* protocol to determine $\max(A^S + B^S, A^U + B^U)$. The system gets the result.

Speaker identification

Input: The system possesses a set of models $\Lambda_1, \Lambda_2, \dots, \Lambda_N$ corresponding to speakers $1 \dots N$ (we consider the “background” model to be one of the speakers). The user has $\mathbf{X} = [\mathbf{x}_1, \dots, \mathbf{x}_T]$.

Output: The system learns $\arg\max_i P(\mathbf{X}; \Lambda_i)$.

- 1) For each speaker $s = 1 \dots N$,
 - a) For each time $t = 1 \dots T$, the user and the system engage in the *SMOG* operation with \mathbf{x}_t and Λ_s to obtain additive shares a_t^s and b_t^s .
 - b) The user computes $A^s = \sum_t a_t^s$. The system computes $B^s = \sum_t b_t^s$.

- 2) The user and system engage in an *SMI* protocol to determine $\arg\max_s A^s + B^s$. The system gets the result.

Speech Recognition

Input: The system possesses a set of models $\Gamma_1, \Gamma_2, \dots, \Gamma_N$ corresponding to words or word sequences $1 \dots N$ (we consider the “background” model to be one of the phrases). The user has $\mathbf{X} = [\mathbf{x}_1, \dots, \mathbf{x}_T]$.

Output: The system learns $\arg\max_i P(\mathbf{X}; \Gamma_i)$. For isolated word recognition or phrase spotting, the speech recognition processes can generally be summarized into the following procedure.

- 1) For each word sequence $s = 1 \dots N$, the user and the system engage in the *SFWD* operation with \mathbf{X} and Γ_s to obtain additive shares A^s and B^s .
- 2) The user and system engage in an *SMV* protocol to determine $\arg\max_s A^s + B^s$. The system gets the result.

The above operation is entirely performed in the log domain – all probabilities are log probabilities – and is hence robust to underflow. In practice the secure Viterbi algorithm is a better choice than the forward algorithm to compute class scores, since its secure implementation has fewer expensive *SMOG* operations, which are replaced by *SMV* operations. In this case $P(\mathbf{X}; \Gamma_s)$ can be computed using the Secure Viterbi Algorithm instead of *SFWD* as described earlier. For continuous speech recognition the optimal state sequence (and the word sequence corresponding to it) can be obtained using the Secure Viterbi Algorithm.

D. Analyzing the protocols

Correctness:

All of the presented protocols, including the primitives, the score computation and the actual classification procedures above are easily shown to be *correct* – the final result obtained with the private protocols is exactly what would have been obtained had the operations been performed in a non-private manner. It can also be ascertained that in practical implementations the insecure and secure versions of the computations are virtually indistinguishable, thus the accuracy tradeoff owing to encryption is negligible [7], [8].

Security:

Each of the operations above is also computationally secure against honest but curious participants. The individual primitives in Section IV-A are easily shown to be secure (e.g. [37]) – Alice and Bob do not learn about

each others' data at all, since they only see ciphertexts, or data masked by additive noise [38]. Consequently, the secure Gaussian operation, the *SMOG* operation, the *SFWD* and the Secure Viterbi (*SVIT*) protocol are all guaranteed not to reveal the user and system's data to one another if the protocols are correctly followed. In general, none of the protocols reveal more than what the outcome of the computation itself reveals.

One might also consider a *malicious* model, where one or more parties may manipulate the data or send bogus data in an attempt to disrupt the protocol or learn about the other party's data. If both parties are malicious, security can be enforced by accompanying the protocols with zero-knowledge proofs [39]. If only one of the parties is malicious, the other party can use conditional disclosure of secrets [40] to make sure he/she receives valid inputs from the malicious party. Both these methods, however, greatly increase the computation and communication overhead of the protocols.

Finally we note that although, technically, the proposed protocols (*SVIT*) also permit fully continuous speech recognition (CSR), they assume that the entire HMM representing the complete word graph to be searched is evaluated for each analysis frame. In practice, CSR is never preformed in this manner, even in conventional non-private implementations. CSR has a high memory footprint and high computational complexity, and the word graphs must be pruned heavily based on partial scores, in order to restrict the computation. The act of pruning restricts the hypothesis set considered and reveals information about the recognition output. Techniques to hide this information are not within the scope of this article, although they are topics of active research.

Performance:

The private computation techniques must in principle result in identical classification outcomes to their conventional non-private counterparts. Although there is a minor loss of resolution resulting from the fact that all computation must now be performed with fixed-point arithmetic to accommodate encryption over integer fields, this has little effect on accuracy – speech processing applications have historically achieved reasonable performance with fixed-point implementations with as little as 16 bits of resolution.

Perhaps the most important performance consideration is the additional computational complexity imposed by the privacy primitives. In particular, computing a single Gaussian securely takes a significant fraction of a second on a desktop computer [7]. Table I reports computation times for 1 second of audio [8] on an isolated word

recognition experiment for a vocabulary of the ten digits 0-10, each of which was modeled by a 5-state HMM with a single Gaussian. Results were obtained on a 3.2 GHz Pentium 4. The Paillier encryption scheme was used [12]. The difference between the forward scores computed using secure computation those computed in the conventional manner was less than 0.52%, and the classification accuracy for the secure and conventional versions was nearly the same, being 99%. Results with different key sizes are given primarily to show the trends in the computation time. It should be noted that Paillier encryption based on 256-bit and 512-bit keys is no longer considered sufficiently secure in the cryptography community.

Table II shows a similar computational expense table containing the average time required to perform privacy-preserving speaker authentication on a Core 2 Duo 2 GHz Linux machine per second of input audio. In this case, BGN encryption [41] is employed, rather than the significantly less expensive Paillier encryption. The data are from the YOHO dataset [42]. Both the speaker and the subject were modeled by mixtures of 32 Gaussians, and the audio was represented by a sequence of 39-dimensional feature vectors computed at the rate of 100 times a second. By comparison, the “insecure” computation on the same experiment took only 3.2 seconds per second of audio. The scores computed by the secure computation were identical to within five decimal places to those obtained in the insecure conventional classifier. Classification accuracies on this data set, with this setup achieves an EER of about 7%. Clearly then, computation is a serious bottleneck. Encryption and decryption take the most time. Additional operations take relatively small time in comparison to the overhead of cipher text processing; nevertheless even they are considerably more expensive than insecure computation. The times shown do not include communication overhead; however for these tasks the communication overhead is trivial compared to the computational expenses.

In the experiments presented above, the implementations were far from optimal. Moreover, the primitives described herein can themselves be optimized; it is not

TABLE I
EXECUTION TIME FOR ISOLATED WORD RECOGNITION PROTOCOL
(PAILLIER ENCRYPTION).

Activity	256-bit keys	512-bit keys	1024-bit keys
Alice encrypts input data (only once)	205.23 s	1944.27 s	11045 s
Bob computes $\xi(\log b_j(\mathbf{x}_t))$ (per HMM)	79.47 s	230.30 s	461 s
Both compute $\xi(\alpha_T(j))$ (per HMM)	16.28 s	107.35 s	785 s

TABLE II
EXECUTION TIME FOR THE VERIFICATION PROTOCOL (PAILLIER
ENCRYPTION).

Steps	Time (256-bit)	Time (1024-bit)
Encrypting $\bar{x}_t \forall t$	138 s	8511 s
Evaluating Adapted	97 s	1809 s
Evaluating UBM	$\approx 97s$	$\approx 1809s$
Comparison	0.07 s	4.01 s
Total $= E[\bar{x}_t] + \text{adapted}$ + UBM + compare	331 s ~ 5.47 min	12133 s ~ 3 hr, 32 min

clear that the particular structure chosen to decompose the computations is optimal from the perspective of computational complexity. Performing computations on parallel processors, more efficient implementations of encryption etc. all may improve the speeds by some orders of magnitude; regardless, the final outcome is currently far slower than the speed performance of insecure computations.

The methods described so far provide privacy to conventional state-of-art mechanisms for performing pattern recognition applications on speech. An alternate mechanism may be to modify the matching algorithms themselves to make them more amenable to efficient secure implementations. In the next section, one such approach is described. These techniques are not as generic in their scope as SMC- based methods described above, since the right form of classifier must be found for the task. Specifically, the technique we describe now applies only to speaker authentication and congruent tasks.

V. SPEECH PROCESSING AS PRIVACY PRESERVING STRING MATCHING

We now discuss an alternative framework for privacy-preserving speech processing based on private string comparison. The main idea is to convert a speech sample into a fingerprint, i.e., a fixed-length bit string. This representation allows us to compare two speech samples by checking if their respective fingerprints match exactly. Unlike the cryptographic framework described earlier, string comparison is *non interactive*, and also much faster than encryption which enables us to perform the privacy-preserving processing very efficiently.

The fingerprint representation is similar to a text-password system. The privacy issues are also similar; the user requires the system to store and compare the passwords in an obfuscated form, so that an adversary cannot observe the original passwords. Furthermore, we require the system to be accurate; just as the password system is able to reject users entering incorrect passwords, the privacy-preserving speech processor should be able to classify speech samples accurately. Finally, the performance of this string comparison-based speech processing approach should be competitive with conventional speech processing methods. Initial software

implementations reveal a tradeoff between a significant increase in speed and a small degradation in the accuracy with respect to classical methods based on, for example, Hidden Markov Models.

The twin objectives of accuracy and privacy are achieved as follows: A data-length independent feature vector (string) is derived from the speech signal such that string comparison implies a nearest-neighbor classification. Second, these vectors are converted into password-like bit strings that are not invertible. We describe these briefly below.

A. Feature Representation: Supervectors

An obvious solution to create a fingerprint is to apply a *cryptographic hash function* $H[\cdot]$, e.g., SHA-256 [43], to the speech input itself. However, direct conversion of audio signals into fingerprint-like patterns is difficult, due the inherent variation in cadence and length of audio recordings. A more robust length- and cadence-invariant representation of the audio is first required.

Campbell, et al. [44] extend the MAP adaptation procedures for GMMs mentioned in Section II-A to construct a *supervector* (SV) to represent each speech sample. A supervector is a characterization of an estimate of the distribution of feature vectors derived from the speech recording. It is obtained by performing maximum *a posteriori* (MAP) adaptation of the UBM over the recording and concatenating the parameters, typically just the means, or means and mixture weights of the adapted model. For instance, given the adapted model $\lambda_s = \{\hat{w}_i^s, \hat{\mu}_i^s, \hat{\Sigma}_i^s\}$ with M Gaussian components, the supervector \mathbf{sv} is given by $(\hat{\mu}_1^s \parallel \hat{\mu}_2^s \parallel \dots \parallel \hat{\mu}_M^s)$.

The supervector is then used as a feature vector instead of the original feature vectors derived from the speech sample. In the speaker authentication task addressed by Campbell et al., authentication verification is performed using a binary support vector machine (SVM) classifier for each user. The SVM is trained on supervectors obtained from enrollment utterances from the user, as instances of one class, and from a collection of impostor recordings as instances of the opposite class. As the classes are usually not separable in the original space, [44] also use a kernel mapping that is shown to achieve higher accuracy.

A related approach is to use k -nearest neighbors trained on supervectors as our classification algorithm. The rationale for this approach is twofold: firstly, k -nearest neighbors allow classification with non-linear decision boundaries with accuracy comparable to SVMs with kernels [45]. Secondly, using the LSH transformations discussed below, private k -nearest neighbors




Fig. 1. Locality Sensitive Hashing

computation can be reduced to private string comparison, which can be easily accomplished without an interactive protocol.

B. Locality Sensitive Hashing

Locality sensitive hashing (LSH) [46] is a widely used technique for performing efficient approximate nearest-neighbor search. An LSH function $L(\cdot)$ proceeds by applying a random transformation to a data vector \mathbf{x} to projecting it to a vector $L(\mathbf{x})$ into a lower dimensional space, which we refer to as the LSH *key* or *bucket*. A set of data points that map to the same key are considered as approximate nearest neighbors (Figure 1).

A single LSH function does not group the data points into fine-grained clusters, one must use a hash key obtained by concatenating the output of k LSH functions. This k -bit LSH function $L(\mathbf{x}) = L_1(\mathbf{x}) \cdots L_k(\mathbf{x})$ maps a d -dimensional vector into a k -bit string. Two data vectors may be deemed to be highly similar if the k -bit hashes derived from them are identical. Such fine selectivity may however have an adverse effect on the recall in identifying neighbors when the data from a class have significant inherent variations. To address this, m different LSH keys are computed over the same input to achieve better recall. Two data vectors \mathbf{x} and \mathbf{y} are said to be neighbors if at least one of their keys, each of length k , matches exactly. LSH provides a major efficiency advantages: By precomputing the keys, the approximate nearest neighbor search can be done in time sub-linear in the size of the dataset.

A family of LSH functions is defined for a particular distance metric. A hash function from this family has the property that data points that are close to each other as defined by the distance metric are mapped to the same key with high probability. There exist LSH constructions for a variety of distance metrics, including arbitrary kernels [47], but we mainly consider LSH for Euclidean distance (E2LSH) [48] and cosine distance [49] as the LSH functions for these constructions are simply random vectors. As the LSH functions are data independent, it is possible to distribute them to multiple parties without privacy loss.

The LSH construction for Euclidean distance transforms a d -dimensional vector into a vector of k integers. The i^{th} entry of the hash is as follows.

$$L_i(\mathbf{x}) = \left\lfloor \frac{\mathbf{r}_i^T \mathbf{x} + b}{w} \right\rfloor, \quad (8)$$

where \mathbf{r}_i is a d -dimensional vector with each component drawn iid from $\mathcal{N}(0, 1)$, w is the width of the bin, (e.g., 255), and $b \in [0, w]$. Similarly, the construction of the i^{th} bit of the LSH for cosine distance, using \mathbf{r}_i defined as above is given by

$$L_i(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{r}_i^T \mathbf{x} > 0, \\ 0 & \text{otherwise,} \end{cases} \quad (9)$$

LSH is inherently not privacy preserving due to its locality sensitive property. It is possible to reconstruct the input vector by observing a sufficient number of LSH keys obtained from the same vector. To satisfy the privacy constraint, a cryptographic hash function $H[\cdot]$ is applied to the LSH keys. Cryptographic hash functions, such as SHA-256, MD5, are orders of magnitude faster to compute compared to homomorphic encryption.

C. Privacy-Preserving Speech Processing through String Comparison

The private string comparison framework lends itself very well to biometric tasks where audio-length-invariant supervector representations of the audio may be derived. In applications such as speaker verification and speaker identification, the user can convert the test speech sample into supervectors, apply the LSH transformation, and finally apply a cryptographic hash function and submit the output to the system. The system can simply compare the hashes provided by the user to the hashes computed over the enrollment data and accept or reject the user by comparing the number of matches to a pre-calibrated threshold. Due to the irreversibility of the hashes, the system is not able to recover the original speech sample, and we are able to maintain the privacy of the speech input. As an illustrative example, we present the string comparison technique applied to the privacy-preserving speaker verification problem below and refer the reader to [10] for other applications.

D. Privacy-Preserving Speaker Verification as String Comparison

As the possible values for LSH keys discussed above lie in a relatively small set by cryptographic standards, 256^k for k -bit Euclidean LSH and 2^k for k -bit cosine LSH, it is possible for the server to obtain $L(\mathbf{s})$ from $H[L(\mathbf{s})]$ by applying brute-force search. To make this




Fig. 2. Speaker Verification as String Comparison

attack infeasible, the domain of the hash function $H[\cdot]$ is increased by concatenating the LSH key with a long random string \mathbf{q}_i (e.g., 80-bits in length) that is unique to the user i , which is called the *salt*, as summarized in Fig. 2. Requiring the user to keep the salt private and unique to each system, also gives the additional advantage of rendering cryptographically hashed enrollment data useless to an adversary. With this modification, LSH-based privacy-preserving speaker verification is achieved using the enrollment and authentication protocols described below:

A. Enrollment Protocol

Each user has a set of enrollment utterances $\{\mathbf{x}_1, \dots, \mathbf{x}_n\}$. The users also obtain the UBM and the l LSH functions $\{L_1(\cdot), \dots, L_l(\cdot)\}$, each of length k -bit from the system. Each user i generates the random 80-bit salt string \mathbf{q}_i .

For each enrollment utterance \mathbf{x}_j , user i :

- performs adaptation of \mathbf{x}_j with the UBM to obtain supervector \mathbf{s}_j .
- applies the l LSH functions to \mathbf{s}_j to obtain the keys $\{L_1(\mathbf{s}_j), \dots, L_l(\mathbf{s}_j)\}$.
- applies the cryptographic hash function salted with \mathbf{q}_i to each of these keys to obtain $\{H[L_1(\mathbf{s}_j) \parallel \mathbf{q}_i], \dots, H[L_l(\mathbf{s}_j) \parallel \mathbf{q}_i]\}$, and sends them to the system.

B. Authentication Protocol

For a test utterance \mathbf{x}' , user i :

- performs adaptation of \mathbf{x}' with the UBM to obtain supervector \mathbf{s}' .
- applies the l LSH functions to \mathbf{s}' to obtain the keys $\{L_1(\mathbf{s}'), \dots, L_l(\mathbf{s}')\}$.
- applies the cryptographic hash function salted with \mathbf{q}_i to each of these keys to obtain $\{H[L_1(\mathbf{s}') \parallel \mathbf{q}_i], \dots, H[L_l(\mathbf{s}') \parallel \mathbf{q}_i]\}$, and sends it to the system.
- The system computes the number of matches between the hashed keys for the test utterance and the corresponding hashes of enrollment utterances. If this number exceeds a threshold, it

TABLE III
AVERAGE EER FOR THE TWO ENROLLMENT DATA CONFIGURATIONS AND THREE LSH STRATEGIES: EUCLIDEAN, COSINE, AND COMBINED (EUCLIDEAN & COSINE).

Enrollment: Only Speaker

Euclidean	Cosine	Combined
15.18%	17.35%	13.80%

Enrollment: Speaker & Imposter

Euclidean	Cosine	Combined
15.16%	18.79%	11.86%

accepts the user.

$$\begin{aligned}
 match &= \sum_{i \in \text{enrollment}} \sum_{j=1}^l I(H[L_j(\mathbf{s}_i) \parallel \mathbf{q}_i] \\
 &= H[L_j(\mathbf{s}') \parallel \mathbf{q}_i]), \\
 &\text{if } match > threshold : \text{accept.}
 \end{aligned}$$

The system never observes any LSH key before a salted cryptographic hash function is applied to it. Apart from the salt, the user does not need to store any speech data on its device. The enrollment and verification protocols, therefore, satisfy the privacy constraints discussed above.

E. Experiments

Experiments examining the accuracy and performance of LSH-based schemes have been reported using the YOHO dataset [42] which comprises of a collection of short utterances, each a sequence of three two-digit numbers, produced by 138 speakers. There are 96 enrollment utterances and 40 test utterances from each speaker. Mel-frequency cepstral coefficient (MFCC) features augmented by differences and double differences are chosen as feature vectors in these experiments [10]. A UBM with 64 Gaussian mixture components is trained on a random subset of the enrollment data belonging to all users. Supervectors are obtained by individually adapting all enrollment and verification utterances to the UBM.

Accuracy

The lowest equal error rate (EER) was achieved by using $l = 200$ instances of LSH functions each of length $k = 20$ for both Euclidean and cosine distances. Table III indicates that LSH for Euclidean distance performs better than LSH for cosine distance, while combining the two classifiers gives the best accuracy. Furthermore, using imposter data achieves

lower EER when using the combined scores for both the distances: one-sided classifiers are clearly not sufficiently accurate. While the actual error rates in this example may appear relatively high, it should be noted that an SVM-based classifier working from supervectors obtained a classification accuracy of 10.8% on the same setup. In other experiments not reported here, using supervectors derived from more detailed GMMs, EERs of approximately 5% have been obtained using the string matching framework.

Execution Time

Compared to a non-private speaker recognition system based on supervectors, the only computational overhead for the privacy-preserving version is in applying the LSH and salted cryptographic hash function. For a $64 \times 39 = 2496$ -dimensional supervector representing a single utterance, the computation for both Euclidean and cosine LSH involves a multiplication with a random matrix of size 20×2496 which requires a fraction of a millisecond. Performing this operation 200 times required 15.8 milliseconds on average [10]. The reported times are for a laptop running 64-bit Ubuntu 11.04 with 2 GHz Intel Core 2 Duo processor and 3 GB RAM.

The Euclidean and cosine LSH keys of length $k = 20$ require 8×20 bits = 20 bytes and 20 bits = 1.6 bytes for storage respectively. Using a C++ implementation of SHA-256 cryptographic hashing algorithm based on the OpenSSL libraries [50], hashing 200 instances of each of these keys in total required 28.34 milliseconds on average. Beyond this, the verification protocol only consists of matching the 256-bit long cryptographically hashed keys derived from the test utterance to those obtained from the enrollment data.

VI. CONCLUSIONS

The two frameworks presented in this article both show promise in enabling privacy-preserving speech processing. The cryptographic framework, while more generic, carries the usual computational overhead of cryptography. However, it can be made arbitrarily secure and flexible. The string-matching framework, on the other hand, is much more efficient; however it is restricted in its applicability and currently results in a degradation of performance.

At this point, both of them can only be viewed as initial forays into the implementation of truly secure speech-processing frameworks, and much work remains. Researchers continue to investigate more efficient protocols, possibly using simpler encryption techniques which could be orders of magnitude faster than the methods described here. Within the string-matching framework,

recent work has shown that significantly greater accuracies can be obtained using nearest-neighbor methods such as those described in [34]. These also show great promise in affording greater generalizability than the string-matching solutions described here.

A computationally efficient implementation of a fully homomorphic encryption (FHE) scheme [27] would significantly change the construction of privacy preserving protocols. This would allow the creation of a non-interactive protocol where the user uploads the encrypted speech sample, and the server can perform all the necessary computations without requiring any further involvement from the user. Such a scheme would not alleviate all the problems of privacy-preserving computation; there would still be a need to deal with key distribution, malicious adversaries, communication overhead, and many other practical problems. Nevertheless, researchers and practitioners of privacy-preserving speech processing are following developments in FHE with interest.

REFERENCES

- [1] US Federal Law, Title 18, Part 1, Chapter 19, Section 2511, "Interception and disclosure of wire, oral, or electronic communications prohibited," January 2012.
- [2] "The end of privacy," *The Economist*, April 1999.
- [3] Paul Schwartz and Joel R. Reidenberg, *Data Privacy Law: A Study of United States Data Protection*, LEXIS law, 1996.
- [4] Helen Nissenbaum, *Privacy in Context - Technology, Policy, and the Integrity of Social Life*, Stanford University Press, 2010.
- [5] R. K. Nichols and P. C. Lekkas, *Speech cryptology*, McGraw-Hill, 2002.
- [6] Madhusudana Shashanka and Paris Smaragdis, "Secure sound classification: Gaussian mixture models," in *ICASSP*, 2006.
- [7] Paris Smaragdis and Madhusudana Shashanka, "A framework for secure speech recognition," *IEEE Transactions on Audio, Speech and Language Processing*, vol. 15, no. 4, pp. 1404–1413, 2007.
- [8] Manas A. Pathak, Shantanu Rane, Wei Sun, and Bhiksha Raj, "Privacy preserving probabilistic inference with hidden Markov models," in *ICASSP*, 2011.
- [9] Manas A. Pathak and Bhiksha Raj, "Privacy preserving speaker verification using adapted GMMs," in *Interspeech*, 2011, pp. 2405–2408.
- [10] Manas A. Pathak, *Privacy-Preserving Machine Learning for Speech Processing*, Ph.D. thesis, Carnegie Mellon University, 2012.
- [11] Ronald L. Rivest, Leonard Adleman, and Michael L. Dertouzos, "On data banks and privacy homomorphisms," in *Foundations of Secure Computation*, 1978, pp. 169–180.
- [12] Pascal Paillier, "Public-key cryptosystems based on composite degree residuosity classes," in *EUROCRYPT*, 1999.
- [13] J. Benaloh, "Dense Probabilistic Encryption," in *Proceedings of the Workshop on Selected Areas of Cryptography*, Kingston, ON, Canada, May 1994, pp. 120–128.
- [14] I. Damgård, M. Geisler, and M. Krøigård, "Homomorphic encryption and secure comparison," *International Journal of Applied Cryptography*, vol. 1, no. 1, pp. 22–31, 2008.
- [15] Andrew Yao, "Protocols for secure computations (extended abstract)," in *IEEE Symposium on Foundations of Computer Science*, 1982.

- [16] Michael Rabin, "How to exchange secrets by oblivious transfer," Tech. Rep. TR-81, Harvard University, 1981.
- [17] Andrew Yao, "How to Generate and Exchange Secrets," in *Proceedings of the 27th Annual Symposium on Foundations of Computer Science (SFCS)*, Washington, DC, USA, 1986, pp. 162–167, IEEE Computer Society.
- [18] Manas A. Pathak and Bhiksha Raj, "Privacy preserving speaker verification as password matching," in *ICASSP*, 2012.
- [19] X. Huang, A. Acero, and H.-W. Hon, *Spoken Language Processing*, Prentice-Hall, 2001.
- [20] S.B.Davis and P. Mermelstein, "Comparison of parametric representations for monosyllabic word recognition in continuously spoken sentences," *IEEE Transactions on Acoustics, Speech, and Signal Processing*, vol. 28, 1980.
- [21] S. Imai, "Cepstral analysis synthesis on the mel frequency scale," in *Proc. IEEE International Conference on Acoustics, Speech and Signal Processing*, 1983, vol. 8, pp. 93–96.
- [22] D. Reynolds and R. C. Rose, "Robust text-independent speaker identification using gaussian mixture speaker models," *IEEE Transactions on Speech and Audio Processing*, vol. 3, 1995.
- [23] P. Kenny, "New map estimators for speaker recognition," in *Proc. Eurospeech*, 2003.
- [24] P. Kenny, "Joint factor analysis of speaker and session variability : Theory and algorithms," Technical report CRIM-06/08-13, 2005.
- [25] T. Petrie, "Probabilistic functions of finite state markov chains," *Annals of Mathematical Statistics*, vol. 40, no. 1, pp. 97–115, 1969.
- [26] B. Raj, R. Singh, and T. Virtanen, "The basics of automatic speech recognition," in *Techniques for noise robustness in automatic speech recognition*, T. Virtanen, R. Singh, and B. Raj, Eds. Wiley, 2012.
- [27] "Recent advances in fully homomorphic encryption: a possible future for signal processing in the encrypted domain," *IEEE Signal Processing Magazine. SI Signal Processing in the Encrypted Domain*, 2013.
- [28] S. Goldwasser and S. Micali, "Probabilistic encryption," *Journal of Computer and System Sciences*, vol. 28, pp. 270–299, 1984.
- [29] Jaideep Vaidya, Chris Clifton, Murat Kantarcioglu, and Scott Patterson, "Privacy-preserving decision trees over vertically partitioned data," *TKDD*, vol. 2, no. 3, 2008.
- [30] W. Du and M. J. Atallah, "Privacy-preserving cooperative statistical analysis," in *Computer Security Applications Conference (ACSAC)*. IEEE, 2001, pp. 102–110.
- [31] Shai Avidan and Moshe Butman, "Blind Vision," in *Proceedings of the 9th European Conference on Computer Vision (ECCV)*, Graz, Austria, May. 2006.
- [32] B. Goethals, S. Laur, H. Lipmaa, and T. Mielikäinen, "On private scalar product computation for privacy-preserving data mining," *International Conference on Information Security and Cryptology (ICISC)*, pp. 23–25, 2004.
- [33] Mikhail Atallah and Jiangtao Li, "Secure outsourcing of sequence comparisons," *International Journal of Information Security*, vol. 4, no. 4, pp. 277–287, 2005.
- [34] S. Rane and P. Boufounos, "Privacy-preserving nearest neighbor methods," *IEEE Signal Processing Magazine, To Appear*, 2012.
- [35] V. Kolesnikov, A. Sadeghi, and T. Schneider, "Improved garbled circuit building blocks and applications to auctions and computing minima," in *Cryptology and Network Security (CANS)*, Kanazawa, Japan, December 2009, pp. 1–20.
- [36] F. Kerschbaum, D. Biswas, and S. de Hoogh, "Performance comparison of secure comparison protocols," in *Proc. International Workshop on Database and Expert System Applications*, 2009, pp. 133–136.
- [37] B. Goethals, S. Laur, H. Lipmaa, and T. Mielikäinen, "On private scalar product computation for privacy-preserving data mining," in *Proc. Intl. Conference on Information Security and Cryptology*, 2004.
- [38] M. Quisquater L. Genelle, E. Prouff, "Thwarting higher-order side channel analysis with additive and multiplicative maskings," in *CHES 2011*, 2011.
- [39] J. Feigenbaum, "Overview of interactive proof systems and zero-knowledge," in *Contemporary Cryptology: The Science of Information Integrity*. IEEE Press, 1992.
- [40] S. Laur and H. Lipmaa, "Additive conditional disclosure of secrets and applications," *Cryptology ePrint Archive*, Report 2005/378, 2005, <http://eprint.iacr.org/>.
- [41] Dan Boneh, Eu-Jin Goh, and Kobbi Nissim, "Evaluating 2-DNF formulas on ciphertexts," in *Theory of Cryptography Conference*, 2005, pp. 325–341.
- [42] Joseph P. Campbell, "Testing with the YOHO CD-ROM voice verification corpus," in *ICASSP*, 1995, pp. 341–344.
- [43] National Institute for Standards and Technology, *FIPS 180-3: Secure Hash Standard*, 2008.
- [44] William M. Campbell, Douglas E. Sturim, Douglas A. Reynolds, and Alex Solomonoff, "SVM based speaker verification using a GMM supervector kernel and NAP variability compensation," in *ICASSP*, 2006.
- [45] Johnny Mariéthoz, Samy Bengio, and Yves Grandvalet, *Kernel Based Text-Independent Speaker Verification*, John Wiley & Sons, 2008.
- [46] Piotr Indyk and Rajeev Motwani, "Approximate nearest neighbors: Towards removing the curse of dimensionality," in *Proceedings of the ACM Symposium on Theory of Computing*, 1998, pp. 604–613.
- [47] Brian Kulis and Kristen Grauman, "Kernelized locality-sensitive hashing for scalable image search," in *IEEE International Conference on Computer Vision*, 2009, pp. 2130–2137.
- [48] Mayur Datar, Nicole Immorlica, Piotr Indyk, and Vahab S. Mirrokni, "Locality-sensitive hashing scheme based on p-stable distributions," in *ACM Symposium on Computational Geometry*, 2004, pp. 253–262.
- [49] Moses Charikar, "Similarity estimation techniques from rounding algorithms," in *ACM Symposium on Theory of Computing*, 2002.
- [50] "OpenSSL," <http://www.openssl.org/docs/crypto/bn.html>, 2010.